

```
In [1]: import os  
os.getcwd()
```

```
Out[1]: 'C:\\Users\\chitt'
```

```
In [3]: import numpy as np  
  
import pandas as pd #excellent for dataset manipulation  
  
# for data visulization  
import matplotlib.pyplot as plt  
  
#stats visualization  
import seaborn as sns  
  
#Labelencoding to convert categorical data into lowlevel language  
from sklearn.preprocessing import LabelEncoder  
  
#scaling data  
from sklearn.preprocessing import StandardScaler  
  
#data partions  
from sklearn.model_selection import train_test_split  
  
#algorithams  
from sklearn.linear_model import LogisticRegression  
  
from sklearn.tree import DecisionTreeClassifier  
  
from sklearn.ensemble import RandomForestClassifier  
  
from xgboost import XGBClassifier  
  
#accuracy confusion matric and classification report  
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report  
  
import warnings  
  
# To ignore all warnings  
warnings.filterwarnings("ignore")
```

```
In [5]: df=pd.read_csv(r"D:\NIT Daily Task\Oct\21st DIBETIC PREDICTION\DIBETIC PREDICTION")
```

```
In [7]: df.head()
```

Out[7]:

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level
0	Female	80.0	0	1	never	25.19	6.6	120.0
1	Female	54.0	0	0	No Info	27.32	6.6	120.0
2	Male	28.0	0	0	never	27.32	5.7	120.0
3	Female	36.0	0	0	current	23.45	5.0	120.0
4	Male	76.0	1	1	current	20.14	4.8	120.0

In [9]: `df.isna().any()`

Out[9]:

gender	False
age	False
hypertension	False
heart_disease	False
smoking_history	False
bmi	False
HbA1c_level	False
blood_glucose_level	False
diabetes	False
dtype:	bool

In [11]: `df.corr(numeric_only=True)`

Out[11]:

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level
age	1.000000	0.251171	0.233354	0.337396	0.101354	0.110672
hypertension	0.251171	1.000000	0.121262	0.147666	0.080939	0.084429
heart_disease	0.233354	0.121262	1.000000	0.061198	0.067589	0.070066
bmi	0.337396	0.147666	0.061198	1.000000	0.082997	0.091261
HbA1c_level	0.101354	0.080939	0.067589	0.082997	1.000000	0.166733
blood_glucose_level	0.110672	0.084429	0.070066	0.091261	0.166733	1.000000
diabetes	0.258008	0.197823	0.171727	0.214357	0.400660	0.400660

In [15]: `df.shape`

Out[15]: (100000, 9)

In [17]:

```

for column in df.columns: # iterating each column in df.columns
    unique_values = df[column].unique() #finding unique values of each column

    #printing unique values
    print('Column "{}" has unique values: {}'.format(column, unique_values))

```

Column "gender" has unique values: ['Female' 'Male' 'Other']
 Column "age" has unique values: [80. 54. 28. 36. 76. 20. 44. 79. 42. 32. 53. 78. 67. 15. 37. 40. 5. 69. 72. 4. 30. 45. 43. 50. 41. 26. 34. 73. 77. 66. 29. 60. 38. 3. 57. 74. 19. 46. 21. 59. 27. 13. 56. 2. 7. 11. 6. 55. 9. 62. 47. 12. 68. 75. 22. 58. 18. 24. 17. 25. 0.08 33. 16. 61. 31. 8. 49. 39. 65. 14. 70. 0.56 48. 51. 71. 0.88 64. 63. 52. 0.16 10. 35. 23. 0.64 1.16 1.64 0.72 1.88 1.32 0.8 1.24 1. 1.8 0.48 1.56 1.08 0.24 1.4 0.4 0.32 1.72 1.48]
 Column "hypertension" has unique values: [0 1]
 Column "heart_disease" has unique values: [1 0]
 Column "smoking_history" has unique values: ['never' 'No Info' 'current' 'former' 'ever' 'not current']
 Column "bmi" has unique values: [25.19 27.32 23.45 ... 59.42 44.39 60.52]
 Column "HbA1c_level" has unique values: [6.6 5.7 5. 4.8 6.5 6.1 6. 5.8 3.5 6.2 4. 4.5 9. 7. 8.8 8.2 7.5 6.8]
 Column "blood_glucose_level" has unique values: [140 80 158 155 85 200 145 100 130 160 126 159 90 260 220 300 280 240]
 Column "diabetes" has unique values: [0 1]

In [19]: `df["smoking_history"].value_counts()` *#Value count of smoking_history parameter*

Out[19]: smoking_history
 No Info 35816
 never 35095
 former 9352
 current 9286
 not current 6447
 ever 4004
 Name: count, dtype: int64

In [21]: `df["smoking_history"].value_counts()/len(df)` *#finding the percentage*

Out[21]: smoking_history
 No Info 0.35816
 never 0.35095
 former 0.09352
 current 0.09286
 not current 0.06447
 ever 0.04004
 Name: count, dtype: float64

In [23]: *# Replaceing No Info columns with pd.NA*
`df['smoking_history'] = df['smoking_history'].replace('No Info', pd.NA)`

Replace missing values with the mode it is string so we are using mode
`mode_value = df['smoking_history'].mode()[0]`
`df['smoking_history'] = df['smoking_history'].fillna(mode_value)` *#filling no inf*

Printing the updated value counts
`print(df['smoking_history'].value_counts())`

```
smoking_history
never          70911
former         9352
current        9286
not current    6447
ever           4004
Name: count, dtype: int64
```

In [25]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   gender                100000 non-null object
1   age                   100000 non-null float64
2   hypertension          100000 non-null int64
3   heart_disease         100000 non-null int64
4   smoking_history       100000 non-null object
5   bmi                   100000 non-null float64
6   HbA1c_level           100000 non-null float64
7   blood_glucose_level   100000 non-null int64
8   diabetes              100000 non-null int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB
```

In [31]: `df.gender.value_counts()`

```
Out[31]: gender
Female    58552
Male      41430
Other       18
Name: count, dtype: int64
```

In [33]: `df.describe()`

```
Out[33]:
```

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level
count	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000
mean	41.885856	0.07485	0.039420	27.320767	5.527507	100.000000
std	22.516840	0.26315	0.194593	6.636783	1.070672	100.000000
min	0.080000	0.00000	0.000000	10.010000	3.500000	100.000000
25%	24.000000	0.00000	0.000000	23.630000	4.800000	100.000000
50%	43.000000	0.00000	0.000000	27.320000	5.800000	100.000000
75%	60.000000	0.00000	0.000000	29.580000	6.200000	100.000000
max	80.000000	1.00000	1.000000	95.690000	9.000000	100.000000

In [35]: `#removing , in bmi parameter`
`df["bmi"] = [float(str(i).replace(",","")) for i in df["bmi"]]`

```
In [37]: #plotting value_counts of diabetes in graphical representation
df['diabetes'].value_counts().plot(kind='barh')


#Xlabel name
plt.xlabel('count')

#ylabel name
plt.ylabel('diabetes')

#title of the plot
plt.title('count of diabetes and Non diabetes')

#invert ylabes to no diabetes on top
plt.gca().invert_yaxis()

#printing the plot
plt.show()
```



```
In [39]: df['diabetes'].value_counts()/len(df) #percentage of 1--diabetes and 2--no diabe
```

```
Out[39]: diabetes
0      0.915
1      0.085
Name: count, dtype: float64
```

```
In [41]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                 100000 non-null object
1   age                    100000 non-null float64
2   hypertension           100000 non-null int64
3   heart_disease          100000 non-null int64
4   smoking_history        100000 non-null object
5   bmi                    100000 non-null float64
6   HbA1c_level            100000 non-null float64
7   blood_glucose_level    100000 non-null int64
8   diabetes                100000 non-null int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB
```

In [45]: `le=LabelEncoder()`

le

Out[45]:  LabelEncoder

LabelEncoder()

In [47]: `Label_encod_columns=['gender','smoking_history'] #selecting columns to apply La`
`df[Label_encod_columns]=df[Label_encod_columns].apply(le.fit_transform) #applyin`

In [49]: `df.head(3)`

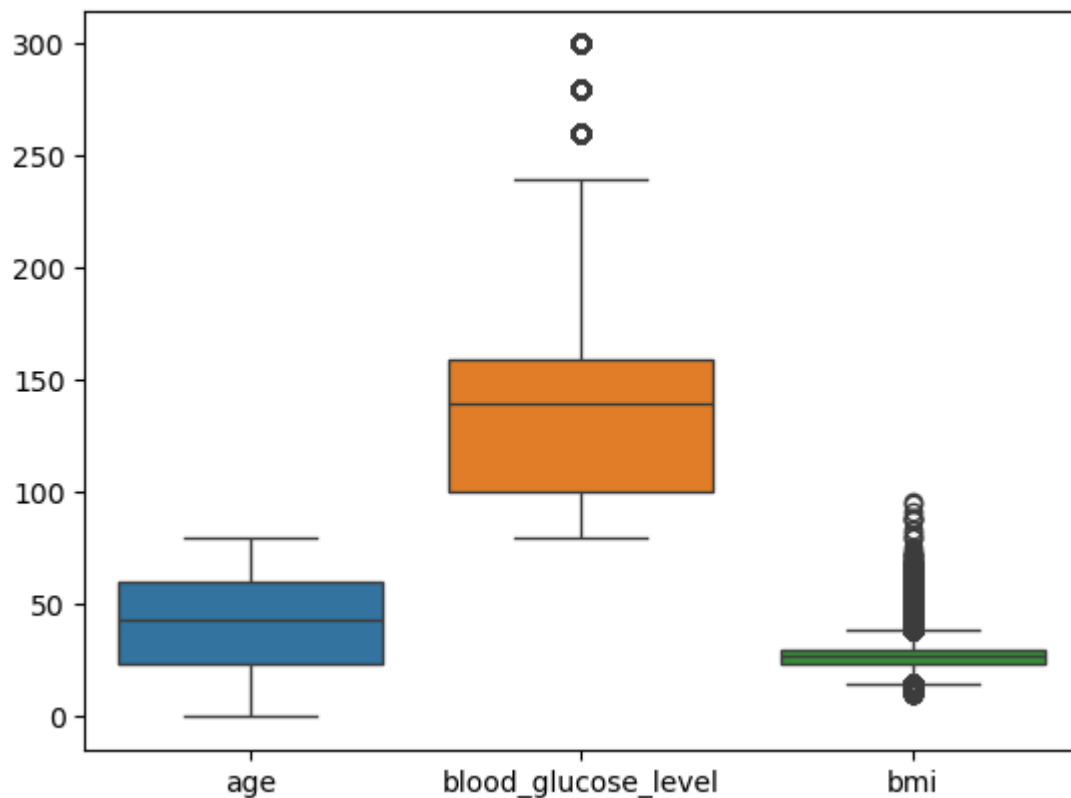
Out[49]:

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blo
0	0	80.0	0	1	3	25.19	6.6	
1	0	54.0	0	0	3	27.32	6.6	
2	1	28.0	0	0	3	27.32	5.7	




In [51]: `sns.boxplot(data=df[['age','blood_glucose_level','bmi']])`

Out[51]: <Axes: >



```
In [53]: sns.boxplot(data=df['HbA1c_level']) #checking outliers using boxplot
```

```
Out[53]: <Axes: ylabel='HbA1c_level'>
```




```
In [55]: ''' it is always good to ignore outliers in medical data '''
```

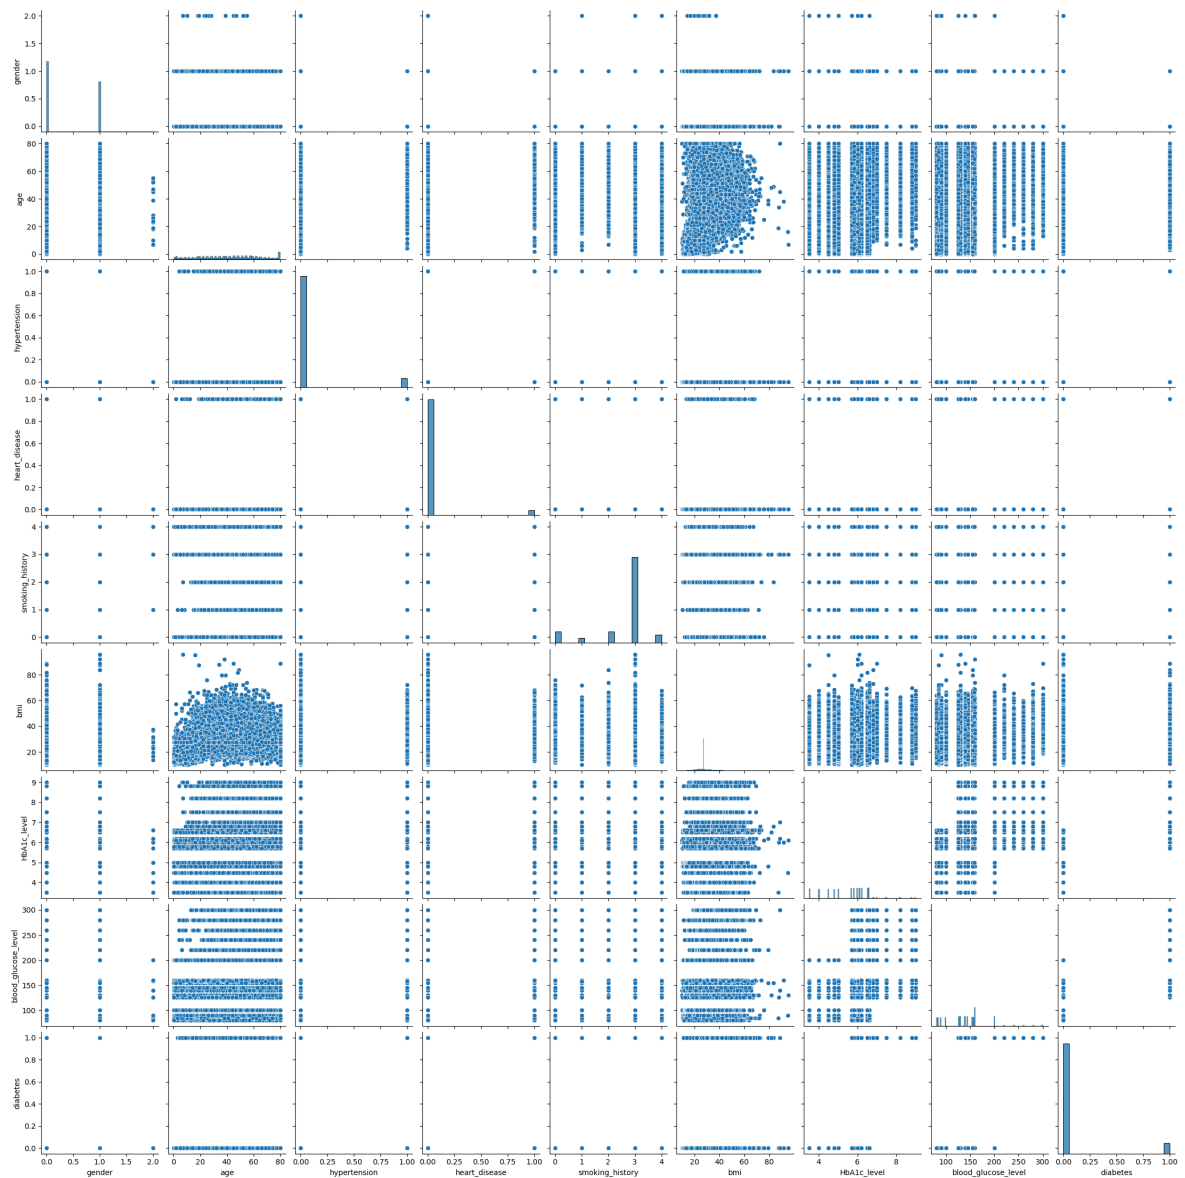
```
Out[55]: ' it is always good to ignore outliers in medical data '
```

```
In [57]: sns.lmplot(data=df, x='blood_glucose_level', y='diabetes', fit_reg=False)#lmplot
```

Out[57]: <seaborn.axisgrid.FacetGrid at 0x24e8d3eb440>



```
In [59]: sns.pairplot(df) #using pairplot to check relation between parameters  
  
         #print the pairplot  
         plt.show()
```

```
In [61]: '''when age increase hypertension and heart disease ,blood_glucose_level and diabetes
relationship between them

*bmi

*HbA1c_level

*blood_glucose_level

these four parameters have relationship between each other

*gender and smoking history it doesnot effect on diabetes

'''
```

```
Out[61]: 'when age increase hypertension and heart disease ,blood_glucose_level and diabetes
and age and also there is a \n relationship between them\n\n *bmi\n\n
\n *HbA1c_level\n\n *blood_glucose_level\n\n these four parameters
have relationship between each other\n\n *gender and smoking history
it doesnot effect on diabetes\n\n'
```


```
In [63]: df.corr()
```

Out[63]:

	gender	age	hypertension	heart_disease	smoking_history
gender	1.000000	-0.030656	0.014203	0.077696	-0.044081
age	-0.030656	1.000000	0.251171	0.233354	-0.098969
hypertension	0.014203	0.251171	1.000000	0.121262	-0.048631
heart_disease	0.077696	0.233354	0.121262	1.000000	-0.048253
smoking_history	-0.044081	-0.098969	-0.048631	-0.048253	1.000000
bmi	-0.022994	0.337396	0.147666	0.061198	-0.087735
HbA1c_level	0.019957	0.101354	0.080939	0.067589	-0.017534
blood_glucose_level	0.017199	0.110672	0.084429	0.070066	-0.022985
diabetes	0.037411	0.258008	0.197823	0.171727	-0.049841

```
In [65]: plt.figure(figsize=(20,8))  
  
df.corr()['diabetes'].sort_values(ascending=False).plot(kind='bar')
```

Out[65]: <Axes: >



In [67]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                100000 non-null  int32
1   age                   100000 non-null  float64
2   hypertension          100000 non-null  int64
3   heart_disease         100000 non-null  int64
4   smoking_history       100000 non-null  int32
5   bmi                   100000 non-null  float64
6   HbA1c_level          100000 non-null  float64
7   blood_glucose_level   100000 non-null  int64
8   diabetes              100000 non-null  int64
dtypes: float64(3), int32(2), int64(4)
memory usage: 6.1 MB
```

```
In [69]: #selecting X variables
X = df.loc[:, 'age':'heart_disease'].join(df.loc[:, 'bmi':'blood_glucose_level'])

X
```

```
Out[69]:
```

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level
0	80.0	0	1	25.19	6.6	140
1	54.0	0	0	27.32	6.6	80
2	28.0	0	0	27.32	5.7	158
3	36.0	0	0	23.45	5.0	155
4	76.0	1	1	20.14	4.8	155
...
99995	80.0	0	0	27.32	6.2	90
99996	2.0	0	0	17.37	6.5	100
99997	66.0	0	0	27.83	5.7	155
99998	24.0	0	0	35.42	4.0	100
99999	57.0	0	0	22.43	6.6	90

100000 rows × 6 columns

```
In [71]: y=df.loc[:, 'diabetes']
y
```

```
Out[71]: 0      0
         1      0
         2      0
         3      0
         4      0
         ..
        99995    0
        99996    0
        99997    0
        99998    0
        99999    0
        Name: diabetes, Length: 100000, dtype: int64
```

```
In [73]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=0)
```

```
In [75]: X_train.head()
```

```
Out[75]:
```

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level
10382	2.0	0	0	16.45	6.2	159
73171	55.0	0	0	24.59	6.0	130
30938	24.0	0	0	21.77	4.5	130
99310	30.0	0	0	27.32	6.2	159
58959	13.0	0	0	18.37	6.5	130

```
In [77]: print('Shape of Train data')
         print(X_train.shape)
         print(y_train.shape)
         print('Shape of Testing data')
         print(X_test.shape)
         print(y_test.shape)
```

```
Shape of Train data
(80000, 6)
(80000,)
Shape of Testing data
(20000, 6)
(20000,)
```

```
In [79]: ss=StandardScaler() #activating StandardScaler()
         ss
```

```
Out[79]: StandardScaler ⓘ ?
         StandardScaler()
```

```
In [81]: X_train_scaled=ss.fit_transform(X_train) #scaling X_train data
```

```
In [83]: if len(X_test.shape) == 1:  #if x is 1d array
        X_test = X_test.values.reshape(-1, 1) #converting to 2d array

        X_test_scaled = ss.fit_transform(X_test) #scaling X_test data
```

```
In [85]: model_lr=LogisticRegression() #activating Logistic Regression
```

```
In [87]: model_lr.fit(X_train_scaled,y_train) #training logistic regression model
```

```
Out[87]: LogisticRegression
```

```
LogisticRegression()
```

```
In [89]: y_pred=model_lr.predict(X_test_scaled) #predecting y_test data
        y_pred[:10]
```

```
Out[89]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int64)
```

```
In [91]: y_test[:10] # actual y_test data
```

```
Out[91]: 3582      0
        60498     0
        53227     0
        21333     0
        3885      0
        51521     0
        84261     0
        10685     1
        59948     0
        41032     0
        Name: diabetes, dtype: int64
```

```
In [95]: accuracy_score(y_pred,y_test)
```

```
Out[95]: 0.95975
```

```
In [97]: print(classification_report(y_pred,y_test)) #classifiaction_report
```

	precision	recall	f1-score	support
0	0.99	0.97	0.98	18736
1	0.63	0.86	0.73	1264
accuracy			0.96	20000
macro avg	0.81	0.91	0.85	20000
weighted avg	0.97	0.96	0.96	20000

```
In [99]: ''' As you can see that the accuracy is quite low, and as it's an imbalanced dat
        Hence, we need to check recall, precision & f1 score for the minority class, and
        Hence, moving ahead to call SMOTEENN (UpSampling + ENN)'''
```

```
'''main advantage of using SMOTEENN is that it addresses both overfitting and un
```

Out[99]: 'main advantage of using SMOTEENN is that it addresses both overfitting and underfitting issues that can arise from class imbalance. By generating synthetic samples and removing noisy ones'


In [101... `confusion_matrix(y_pred,y_test)`

Out[101... `array([[18114, 622],
[183, 1081]], dtype=int64)`

In [103... `y_train.value_counts()`


Out[103... `diabetes`
 0 73203
 1 6797
 Name: count, dtype: int64

In [105... `value_counts=y_train.value_counts()`
`plt.figure(figsize=(16, 8))`
`plt.pie(value_counts, labels=value_counts.index, autopct='%1.2f%%', startangle=1`
`plt.title('Distribution of y_train')`
`plt.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.`
`plt.show()`



In [107... `from imblearn.over_sampling import SMOTE # using smote function to balance our s`
`smote=SMOTE()`
`X_ovs,y_ovs=smote.fit_resample(X,y) #passing X and y variables to it to balance`
`fig, oversp = plt.subplots()`
`oversp.pie(y_ovs.value_counts(), autopct='%1.2f')`
`oversp.set_title("Over-sampling")`
`plt.show()`

Over-sampling



```
In [109... Xr_train,Xr_test,yr_train,yr_test=train_test_split(X_ovs,y_ovs,train_size=0.7,ra
```

```
In [111... print('train data shape')
print(Xr_train.shape)
print(yr_train.shape)
print('test data shape')
print(Xr_test.shape)
print(yr_test.shape)
```

```
train data shape
(128099, 6)
(128099,)
test data shape
(54901, 6)
(54901,)
```

```
In [113... print('y_train and y_test value_count')
print(yr_train.value_counts())
print(yr_test.value_counts())
```

```
y_train and y_test value_count
diabetes
0    64131
1    63968
Name: count, dtype: int64
diabetes
1    27532
0    27369
Name: count, dtype: int64
```

In [115... `ss=StandardScaler()`

`ss`

Out[115... `StandardScaler`

`StandardScaler()`

In [117... `data=Xr_train,Xr_test`

`xr_train_sc=ss.fit_transform(Xr_train) # scaling our resampling data xr train`

`Xr_test_sc=ss.fit_transform(Xr_test) # scaling our resamplig xr_test data`

In [119... `Xr_train_scaled = pd.DataFrame(xr_train_sc) #Xr_train_scaled converting into the`

`print(Xr_train_scaled.shape)`

`Xr_train_scaled.head()`

`print(yr_train.shape)`

`(128099, 6)`

`(128099,)`

In [121... `Xr_test_scaled=pd.DataFrame(Xr_test_sc) #Xr_test converting into the dataframe`

`print(Xr_test_scaled.shape)`

`Xr_test_scaled.head()`

`(54901, 6)`

Out[121...

	0	1	2	3	4	5
0	0.648387	-0.293492	-0.206627	0.841287	-0.245012	-0.653366
1	-1.095183	-0.293492	-0.206627	-0.404988	0.369655	-0.057241
2	-1.466964	-0.293492	-0.206627	-0.288104	0.369655	-1.459888
3	-0.769875	3.407246	-0.206627	0.288161	0.369655	-1.372222
4	-1.374018	-0.293492	-0.206627	-0.288104	-2.152519	-1.109226

In [123... `model_lk=LogisticRegression()`

`model_lk.fit(Xr_train_scaled,yr_train) #trining the model`

Out[123... `LogisticRegression`

`LogisticRegression()`

In [125... `y_pred_lr=model_lk.predict(Xr_test_scaled) #predecting yr_test data`

`y_pred_lr[:10]`

Out[125... array([0, 0, 0, 0, 0, 1, 0, 0, 0, 1], dtype=int64)

In [127... yr_test[:10]

Out[127... 180328 1
573 0
13494 0
93981 0
75389 0
180973 1
71021 0
19293 0
16393 0
121419 1
Name: diabetes, dtype: int64

In [129... *#classification_report for predict value and original value*

print(classification_report(y_pred_lr,yr_test))

	precision	recall	f1-score	support
0	0.88	0.88	0.88	27317
1	0.89	0.88	0.88	27584
accuracy			0.88	54901
macro avg	0.88	0.88	0.88	54901
weighted avg	0.88	0.88	0.88	54901

In [131... *#confusion_matrix for predict value and original value*

confusion_matrix(y_pred_lr,yr_test)

Out[131... array([[24165, 3152],
[3204, 24380]], dtype=int64)

Decision Tree Classifier

In [134... *# activating DecisionTree Classifier*
model_dtc=DecisionTreeClassifier()

passing xr_train_scaled, yr_train to training the model
model_dtc.fit(Xr_train_scaled,yr_train)

model_dtc

Out[134... ▼ DecisionTreeClassifier ⓘ ?
DecisionTreeClassifier()

In [136... y_pred_dtc=model_dtc.predict(Xr_test_scaled) *# predicting yr_test data*

In [138... *# classification report for decisionTreeClassifier*

print(classification_report(y_pred_dtc,yr_test))

	precision	recall	f1-score	support
0	0.79	0.99	0.88	21954
1	0.99	0.83	0.90	32947
accuracy			0.89	54901
macro avg	0.89	0.91	0.89	54901
weighted avg	0.91	0.89	0.89	54901

In [140...] `confusion_matrix(y_pred_dtc,yr_test)`

Out[140...] `array([[21675, 279],
[5694, 27253]], dtype=int64)`

In [142...] `model_rfc=RandomForestClassifier() #activating the fuction`
`model_rfc.fit(Xr_train_scaled,yr_train)`

Out[142...] `RandomForestClassifier`
`RandomForestClassifier()`

In [144...] `y_pred_rfc=model_rfc.predict(Xr_test_scaled)`

In [146...] `print(classification_report(y_pred_rfc,yr_test))`

	precision	recall	f1-score	support
0	0.92	0.99	0.95	25424
1	0.99	0.92	0.95	29477
accuracy			0.95	54901
macro avg	0.95	0.95	0.95	54901
weighted avg	0.95	0.95	0.95	54901

In [150...] `confusion_matrix(y_pred_rfc,yr_test)`

Out[150...] `array([[25054, 370],
[2315, 27162]], dtype=int64)`

XGBOOST

In [154...] `model_xgb=XGBClassifier()`
`model_xgb.fit(Xr_train_scaled,yr_train)`

Out[154...

XGBClassifier

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
```

In [156...

```
y_pred_xgb=model_xgb.predict(Xr_test_scaled)
```

In [158...

```
print(classification_report(y_pred_xgb,yr_test))
```

	precision	recall	f1-score	support
0	0.76	0.98	0.86	21230
1	0.99	0.81	0.89	33671
accuracy			0.88	54901
macro avg	0.88	0.90	0.87	54901
weighted avg	0.90	0.88	0.88	54901

In [160...

```
confusion_matrix(y_pred_xgb,yr_test)
```

Out[160...

```
array([[20896,  334],
       [ 6473, 27198]], dtype=int64)
```

finding the hyperparameter tuning and best param grid

In [163...

```
from sklearn.model_selection import GridSearchCV, cross_val_score
from sklearn.linear_model import LogisticRegression

# Define the parameter grid to search over
param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100], # Regularization parameter
    'penalty': ['l1', 'l2'] # Penalty type
}

# Create a Logistic Regression model
logistic = LogisticRegression()

# Create a GridSearchCV object
grid_search = GridSearchCV(estimator=logistic, param_grid=param_grid, cv=10)

# Initialize an empty list to store the accuracy scores
accuracy_scores = []

# Perform cross-validation 10 times
for _ in range(10):
```

```

# Fit the GridSearchCV object to the training data
grid_search.fit(Xr_train_scaled, yr_train)

# Get the best parameters
best_params = grid_search.best_params_

# Perform cross-validation with the best model
cv_scores = cross_val_score(grid_search.best_estimator_, Xr_train_scaled, yr
# Store the mean accuracy score
accuracy_scores.append(cv_scores.mean())

# Print the accuracy scores obtained over 10 iterations
#print("Accuracy scores over 10 iterations:", accuracy_scores)
print("Accuracy scores over 10 iterations:", ["{:.2f}".format(score) for score in accuracy_scores])

# Get the best parameters and best score
best_params = grid_search.best_params_
best_score = grid_search.best_score_

print("Best parameters found:", best_params)
print("Best cross-validation score:", best_score)

```

Accuracy scores over 10 iterations: ['0.88', '0.88', '0.88', '0.88', '0.88', '0.88', '0.88', '0.88', '0.88', '0.88']
 Best parameters found: {'C': 0.001, 'penalty': 'l2'}
 Best cross-validation score: 0.8848390832485136

Final Model

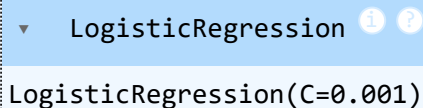
```

In [166... from sklearn.linear_model import LogisticRegression

# Create a Logistic Regression model with the best parameters
final_model = LogisticRegression(C=0.001, penalty='l2')

# Fit the final model to the entire training dataset
final_model.fit(Xr_train_scaled, yr_train)

```

Out[166...  LogisticRegression(C=0.001)

```

In [168... import pickle

# Save the final model to a pickle file
with open('final_model.pkl', 'wb') as file:
    pickle.dump(final_model, file)

```

```

In [170... import pickle
import numpy as np

# Load the model from the pickle file
with open('final_model.pkl', 'rb') as file:
    loaded_model = pickle.load(file)

# Define the mean and standard deviation of the training data
mean_values = [41.885856, 0.07485, 0.03942, 27.320767, 5.527507, 138.058060]

```

```
std_values = [22.516840, 0.26315, 0.194593, 6.636783, 1.070672, 40.708136]

# Define the input features for prediction
age = 30
hypertension = 0
heart_disease = 0
bmi = 100.0
HbA1c_level = 5.0
blood_glucose_level = 90

# Scale the input features manually
scaled_features = [(x - mean) / std for x, mean, std in zip(
    [age, hypertension, heart_disease, bmi, HbA1c_level, blood_glucose_level],
    mean_values, std_values
)]

# Make predictions on the scaled data
prediction = loaded_model.predict([scaled_features])

# Print the prediction
if prediction[0] == 1:
    print("Diabetic")
else:
    print("Not Diabetic")
```

Diabetic

Completed

In []: